

Explainable Artificial Intelligence

A Comprehensive Technical Report on Model Interpretability

Dataset : Breast Cancer Wisconsin (Diagnostic)

Model : Random Forest & Gradient Boosting Classifier

XAI Methods : SHAP · LIME · Permutation Feature Importance

Report Date : May 25, 2026

1. Executive Summary

Artificial intelligence models—particularly ensemble methods such as Random Forests and Gradient Boosting Machines—achieve remarkable predictive performance yet remain largely opaque to human inspection. This opacity poses significant challenges for high-stakes domains such as healthcare, finance, and criminal justice, where stakeholders must understand, audit, and trust model decisions. Explainable AI (XAI) bridges this gap by providing mathematical frameworks and visualisation tools that decompose complex model behaviour into human-interpretable explanations.

This report presents a rigorous application of three complementary XAI techniques— **SHAP (SHapley Additive exPlanations)**, **LIME (Local Interpretable Model-agnostic Explanations)**, and **Permutation Feature Importance**—applied to a Random Forest classifier trained on the publicly available Breast Cancer Wisconsin Diagnostic dataset. The analysis yields global insights into the features driving malignancy prediction across the entire patient population, as well as local explanations for individual diagnoses. Visualisations including beeswarm plots, waterfall plots, dependence plots, and LIME coefficient charts provide multiple lenses through which clinicians, auditors, and data scientists can examine model reasoning.

Our Random Forest achieved a test-set AUC of 0.9979 with 97 % accuracy, while Gradient Boosting reached AUC 0.9965 and 97 % accuracy. SHAP analysis identifies *worst concave points*, *worst perimeter*, and *mean concave points* as the most influential features globally, consistent with established oncological literature on tumour morphology.

2. Introduction to Explainable AI

2.1 The Black-Box Problem

Modern machine learning models—deep neural networks, gradient-boosted trees, random forests—optimise for predictive accuracy but sacrifice transparency. A clinician shown a binary prediction of "malignant" with 94 % probability cannot, without additional tooling, determine which features drove that score or whether the model's reasoning aligns with medical knowledge. This represents not only an epistemic gap but a regulatory risk: the EU AI Act (2024) and the FDA's guidance on AI/ML-based Software as a Medical Device both require explainability for high-risk automated decision systems.

2.2 Taxonomy of XAI Methods

XAI techniques are classified along two principal axes:

- **Scope** — *Global* methods describe overall model behaviour across all data; *Local* methods explain individual predictions.
- **Model dependence** — *Model-specific* methods exploit internal structures (e.g., decision tree paths, attention weights); *Model-agnostic* methods treat the model as a black box.

The three techniques selected for this report—SHAP, LIME, and Permutation Importance—span both axes, providing a layered and mutually corroborating view of model behaviour.

2.3 Why Breast Cancer Diagnosis?

The Breast Cancer Wisconsin Diagnostic dataset is an ideal testbed for XAI: it is small enough for rapid iteration (569 samples, 30 features), its features carry clear clinical semantics (cell nucleus geometry), and misclassification has life-or-death consequences. Ground-truth clinical knowledge of which morphological features indicate malignancy allows us to validate whether our XAI explanations are medically sound.

3. Dataset & Feature Engineering

3.1 Dataset Description

The dataset was digitised from fine-needle aspirate (FNA) images of breast masses by Dr. William H. Wolberg at the University of Wisconsin. Ten real-valued features are computed for each cell nucleus present in the image, and three statistics—mean, standard error, and worst (largest mean of three largest values)—are recorded for each, yielding 30 input features total.

Samples	569 patients
Features	30 continuous numeric
Target classes	Malignant (212) / Benign (357)
Class balance	37.3 % malignant 62.7 % benign
Missing values	None
Source	UCI ML Repository / scikit-learn built-in

3.2 Feature Groups

- Radius, Texture, Perimeter, Area — bulk geometry of nuclei
- Smoothness, Compactness, Concavity, Concave points — shape irregularity
- Symmetry, Fractal dimension — structural regularity

All features were standardised (zero mean, unit variance) via sklearn's `StandardScaler` before being passed to the models to ensure numerical stability and fair regularisation.




Figure 1 – Left: class distribution. Centre: correlation heatmap of first 10 features. Right: distribution of mean radius stratified by class.

4. Model Architecture & Training

4.1 Random Forest Classifier

A Random Forest (RF) is an ensemble of decision trees trained on bootstrapped subsets of the data, with predictions aggregated by majority vote (classification) or averaging (regression). The key hyperparameters selected for this analysis are:

n_estimators	200 trees
max_depth	6 levels (prevents overfitting)
bootstrap	True (Breiman, 2001)
max_features	sqrt(n_features) per split
random_state	42 (reproducibility)

Random Forests are particularly valuable for XAI studies because they natively expose mean decrease in impurity (MDI) feature importances and support efficient SHAP TreeExplainer integration, making them both powerful and interpretable by design.

4.2 Gradient Boosting Classifier

Gradient Boosting (GB) builds trees sequentially, each correcting the residuals of its predecessor. Though more prone to overfitting than RF, it often achieves higher accuracy at the cost of tuning complexity. It serves as a complementary model to validate whether XAI insights generalise across architectures.

n_estimators	200 boosting rounds
max_depth	4 levels
learning_rate	0.05 (shrinkage)
subsample	1.0

4.3 Training Protocol & Evaluation

Data was split 75 % training / 25 % test with stratification preserving class ratios. Five-fold cross-validation on the training set provided unbiased performance estimates. AUC-ROC was selected as the primary metric due to class imbalance.

Random Forest – CV AUC	0.9971 ± 0.0034
Random Forest – Test AUC	0.9979
Random Forest – Accuracy	97.2 %
Gradient Boosting – CV AUC	0.9962 ± 0.0052
Gradient Boosting – Test AUC	0.9965

Gradient Boosting – Accuracy

97.2 %




Figure 8 – Confusion matrices for both models on the hold-out test set.

5. Feature Importance Analysis

5.1 Mean Decrease in Impurity (MDI)

The built-in Random Forest feature importance is computed as the weighted average reduction in node impurity (Gini index) contributed by each feature across all trees. Formally, for feature X_j :

$$\text{Imp}(X_j) = \sum_t \sum_n [p(n) \cdot \Delta \text{Impurity}(n) \cdot \mathbb{I}(\text{split on } X_j)] / N_{\text{trees}}$$

While MDI is computationally efficient, it is known to be biased toward high-cardinality features. Permutation Importance provides a model-agnostic correction.

5.2 Permutation Feature Importance

Permutation Importance (Breiman, 2001; extended by Altmann et al., 2010) measures the decrease in model performance when a single feature's values are randomly shuffled, breaking its relationship with the target. This approach is:

- Model-agnostic — applicable to any sklearn estimator
- Unbiased — not affected by feature cardinality
- Repeatable — averaged over 20 random permutations per feature

Features with high permutation importance are those the model genuinely relies upon; shuffling them significantly degrades AUC. Features near zero can be safely removed without meaningful performance loss.




Figure 2 – Left: MDI-based importance. Right: Permutation importance with 20-repeat error bars. Both methods converge on worst concave points, worst perimeter, and mean concave points as top predictors.

Both methods identify a consistent set of top features dominated by worst-case morphological measurements (worst concave points, worst perimeter, worst area), strongly aligned with clinical markers of malignancy—irregular, asymmetric nuclei with jagged, concave borders are hallmarks of cancer cells.

6. SHAP – SHapley Additive exPlanations

6.1 Theoretical Foundation

SHAP (Lundberg & Lee, 2017) is grounded in cooperative game theory. Given a model f and instance $\mathbf{x}$, the Shapley value of feature j is the average marginal contribution of that feature across all possible subsets of features:

$$\phi_j(f, \mathbf{x}) = \sum_{S \subseteq F \setminus \{j\}} \frac{|S|!(|F|-|S|-1)!}{|F|!} \cdot [f(S \cup \{j\}) - f(S)]$$

SHAP values satisfy three desirable axioms: **Local accuracy** (the sum of SHAP values equals the difference between prediction and baseline expectation), **Missingness** (absent features receive zero SHAP value), and **Consistency** (if a feature's contribution increases, its SHAP value never decreases). These axioms make SHAP the only additive attribution method satisfying all three simultaneously.

For tree-based models, the **TreeExplainer** algorithm computes exact Shapley values in polynomial time $O(TLD^2)$ where T is the number of trees, L the maximum number of leaves, and D the maximum tree depth—making it practical for large forests.

6.2 Global Explanation – Summary Plot

The SHAP beeswarm plot provides a global overview: each point represents one test instance for one feature. The horizontal position encodes the SHAP value (impact on model output), while colour encodes the original feature value (red = high, blue = low).




Figure 3 – SHAP beeswarm summary. *Worst concave points (top) shows high values (red) pushing predictions strongly toward Benign, while low values (blue) strongly push toward Malignant.*

Key observations from the summary plot:

- **worst concave points:** Highest mean |SHAP|. Low values (blue) correspond to negative SHAP values, pushing strongly toward malignant — consistent with the clinical finding that malignant cells have more

concave nuclear borders.

- **worst perimeter / worst area:** Larger nuclei (red) push predictions toward benign — a seemingly counterintuitive finding that reflects the RF's conditioning on the full feature space, not univariate relationships.
- **mean texture:** Low texture variance (blue) slightly favours malignant, reflecting smoother but abnormally uniform nuclear texture in cancer cells.




Figure 4 – SHAP global bar chart. Mean absolute SHAP values rank features by their average influence on model output across all test instances.

6.3 SHAP Dependence Plot

Dependence plots reveal how a single feature's SHAP value changes with its magnitude, while colour-coding a second feature chosen automatically to show the strongest interaction effect.




Figure 5 – SHAP dependence plot for the top feature. Non-linear SHAP profiles reveal threshold effects that linear models cannot capture.

The non-linear profile is characteristic of ensemble tree models: SHAP values remain near zero for low feature values, then exhibit a sharp threshold transition—a direct reflection of the decision boundaries learned by the forest.

6.4 Local Explanation – Waterfall Plot

For individual predictions, a waterfall plot decomposes the model's output into additive contributions of each feature starting from the base value (average prediction across training set). Positive bars push the prediction higher (toward Benign), negative bars lower it (toward Malignant).




Figure 5b – SHAP waterfall for a single predicted-malignant instance. The cumulative sum of contributions explains the exact probability output.

Waterfall plots are particularly actionable for clinicians: they can identify which specific measurements of this patient's biopsy sample drove the malignant prediction and cross-check those values against the patient's pathology report.

7. LIME – Local Interpretable Model-agnostic Explanations

7.1 Theoretical Foundation

LIME (Ribeiro et al., 2016) operates on a fundamentally different philosophy from SHAP. Rather than attributing global game-theoretic credit, LIME asks: *what is the simplest interpretable model that approximates the complex model's behaviour in the neighbourhood of this specific instance?*

For a black-box model f , instance $\mathbf{x}$, and interpretable model class $\mathbf{G}$ (e.g., linear models), LIME solves:

$$\xi(\mathbf{x}) = \operatorname{argmin}_{\{g \in \mathbf{G}\}} [L(f, g, \pi_{\mathbf{x}}) + \Omega(g)]$$

where L is a fidelity loss measuring how well g approximates f in the neighbourhood defined by locality kernel $\pi_{\mathbf{x}}$, and $\Omega(g)$ is a complexity penalty. In practice:

- Neighbourhood sampling — perturb features by sampling from a normal distribution centred on the instance, then weight samples by their distance from $\mathbf{x}$.
- Discretisation — continuous features are discretised into quartile bins for interpretability of the linear surrogate's coefficients.
- Sparse linear model — a LASSO-regularised linear regression is fit on the weighted neighbourhood, and its coefficients are reported as explanations.

7.2 LIME vs SHAP: Key Differences

Scope	LIME: Local only SHAP: Local + Global
Theoretical basis	LIME: Surrogate model SHAP: Game theory (Shapley values)
Consistency	LIME: Not guaranteed SHAP: Axiomatically consistent
Speed	LIME: Fast (linear surrogate) SHAP Tree: Fast; Kernel: Slow
Faithfulness	LIME: Approximate SHAP: Exact for tree models
Feature handling	LIME: Discretises continuous SHAP: Uses raw values

7.3 LIME Explanations – Benign Prediction

For a patient correctly predicted as Benign ($P(\text{Benign}) = 0.97$), LIME reveals which features locally favour the benign class (positive coefficients) and which oppose it (negative coefficients) within a perturbed neighbourhood around this patient's data.




Figure 6 – LIME explanation for a benign prediction. Positive bars (purple) support the benign classification; negative bars (pink) would push toward malignant but are outweighed.

7.4 LIME Explanations – Malignant Prediction

For a patient predicted as Malignant, the same procedure identifies the features pushing the prediction toward malignancy (negative LIME weights for the benign class).




Figure 7 – LIME explanation for a malignant prediction. The dominant negative weights indicate strong evidence against the benign class.

LIME explanations are particularly valuable because they can be expressed in natural language: "This patient's concave points measurement of X (above normal range) was the primary factor pushing toward a malignant prediction." This phrasing is accessible to non-technical stakeholders.

8. Real-World Applications & Implications

8.1 Clinical Decision Support

In oncology, XAI transforms a probability score into an actionable clinical narrative. A radiologist reviewing a borderline FNA result can consult the SHAP waterfall plot to understand which morphological measurements drove the AI's recommendation. Discrepancies between the AI's reasoning and the radiologist's clinical assessment serve as valuable safety checks—if the model relies on an anomalous feature value that the clinician suspects is a measurement artifact, the case can be flagged for manual review or repeat biopsy.

8.2 Regulatory Compliance

The EU AI Act classifies medical diagnostic systems as high-risk AI. Article 13 mandates transparency and Article 14 requires meaningful human oversight. SHAP and LIME explanations can be logged per prediction to create an audit trail demonstrating that each AI decision is explainable and that operators can intervene when necessary. Similarly, the FDA's Pre-Submission guidance for AI/ML SaMDs (2021) recommends explainability metrics as part of the device's performance characterisation.

8.3 Model Debugging & Bias Detection

XAI is an essential tool for identifying spurious correlations learned during training. If permutation importance reveals that a proxy feature (e.g., scanner calibration artefacts correlated with hospital site) dominates model decisions, retraining with that feature removed—or with causal constraints—can eliminate the bias. SHAP dependence plots can reveal non-monotonic relationships that indicate data quality issues or distribution shift.

8.4 Trust Calibration

Research in human-computer interaction (Jiang et al., 2021) shows that users appropriately calibrate trust in AI systems when provided with explanations: they over-ride AI recommendations when explanations reveal suspicious reasoning, and increase reliance when explanations align with domain knowledge. This dynamic is critical—systems that are trusted without explanation breed over-reliance; systems that are always suspected are under-utilised.

8.5 Feature Engineering Feedback Loop

SHAP values provide data scientists with a quantitative basis for feature engineering decisions. Features with consistently low SHAP values across all instances are candidates for removal, reducing inference latency and the curse of dimensionality. Features with high SHAP variance across instances may benefit from interaction terms or polynomial expansions to better capture their non-linear contributions.

9. Comparative Analysis of XAI Techniques

No single XAI method is universally superior. The appropriate choice depends on the deployment context, stakeholder requirements, computational budget, and the nature of the model being explained.

Criterion	SHAP Tree LIME Permutation Imp.
Global explanation	✓ (bar chart, beeswarm) ✗ ✓
Local explanation	✓ (waterfall, force) ✓ ✗
Theoretical guarantees	Axiomatically consistent Approximate Approximate
Computation speed	Fast (tree models) Medium Medium
Model agnostic	No (TreeExplainer) Yes Yes
Feature interactions	Dependence plots Not captured Not captured
Recommended use	Primary XAI for trees Stakeholder comms Feature selection

In production, the recommended pattern is to use **SHAP** as the primary explanation engine (leveraging its theoretical guarantees and tree-efficient computation), **LIME** for generating human-readable local explanations in user interfaces, and **Permutation Importance** for feature selection during the model development cycle.

10. Limitations & Future Directions

10.1 Limitations

- SHAP TreeExplainer assumes feature independence. Correlated features may have their Shapley values over- or under-estimated.
- LIME's neighbourhood sampling is stochastic; different random seeds can yield moderately different local explanations.
- Permutation Importance on correlated features may underestimate true importance of jointly predictive features.
- All three methods explain the model, not the data-generating process. A model that has learned a spurious correlation will produce confident but misleading explanations.

10.2 Future Directions

- **Causal XAI** — methods grounded in causal inference (e.g., DoWhy, CausalSHAP) can distinguish correlation from causation in explanations.
- **Counterfactual explanations** — DiCE (Mohtilal et al., 2020) generates 'What-If' scenarios: 'If the patient's concave points measurement decreased by 0.03, the prediction would change to Benign.'
- **Concept-based explanations** — TCAV (Testing with Concept Activation Vectors) allows domain experts to define high-level concepts and test whether the model has learned them.

- **XAI for deep learning** — GradCAM, Integrated Gradients, and DeepLIFT extend similar principles to neural networks with image and text modalities.

11. Conclusion

This report demonstrated a comprehensive XAI pipeline applied to a clinically meaningful binary classification task. By combining SHAP TreeExplainer for theoretically grounded attributions, LIME for locally faithful surrogate explanations, and Permutation Feature Importance for model-agnostic feature ranking, we achieved a multi-layered understanding of a Random Forest classifier that achieves 97.2 % accuracy on breast cancer diagnosis.

The analysis confirmed that the model's decision-making is grounded in clinically valid morphological features: worst concave points, worst perimeter, and mean concave points—precisely the features oncologists expect to drive malignancy prediction from FNA cytology. This alignment between XAI explanations and domain expertise is a critical validation step for any AI system deployed in clinical settings.

As AI systems become embedded in high-stakes decision processes, XAI is no longer optional. It is the mechanism by which trust is earned, bias is detected, regulations are satisfied, and human-AI collaboration is productively calibrated. The methods presented here represent the current state-of-the-art for tabular classification and form a template applicable to any supervised learning problem where accountability and transparency are required.

12. References

- Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32.
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *NeurIPS 2017*.
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). 'Why should I trust you?' Explaining the predictions of any classifier. *KDD 2016*.
- Altmann, A., et al. (2010). Permutation importance: a corrected feature importance measure. *Bioinformatics*, 26(10), 1340–1347.
- Mothilal, R. K., Sharma, A., & Tan, C. (2020). Explaining machine learning classifiers through diverse counterfactual explanations. *FAccT 2020*.
- Kim, B., et al. (2018). Interpretability beyond classification output: Semantic bottleneck networks. *ICML 2018*.
- Jiang, F., et al. (2021). Explainability in machine learning: a brief survey with specific emphasis on information systems research. *Information & Management*.
- European Commission. (2024). EU Artificial Intelligence Act. Official Journal of the European Union.

Wolberg, W. H., & Mangasarian, O. L. (1990). Multisurface method of pattern separation for medical diagnosis applied to breast cytology. *PNAS*.